

# Introduction to Generative Adversarial Networks (GANs)


Nguyen Tuan Duong  
Nguyen Phuoc Tat Dat  
(2017/9/9)

# Agenda

- ❖ Generative Modeling
- ❖ Idea of GAN
- ❖ Training GAN
- ❖ Distribution matching
- ❖ Example
- ❖ Variants
- ❖ Applications of Generative Models

# Generative Modeling

# Generative Modeling




Training samples




$p_{\text{data}}$

Density estimation  
(Explicit or implicit)



$p_{\text{model}}$

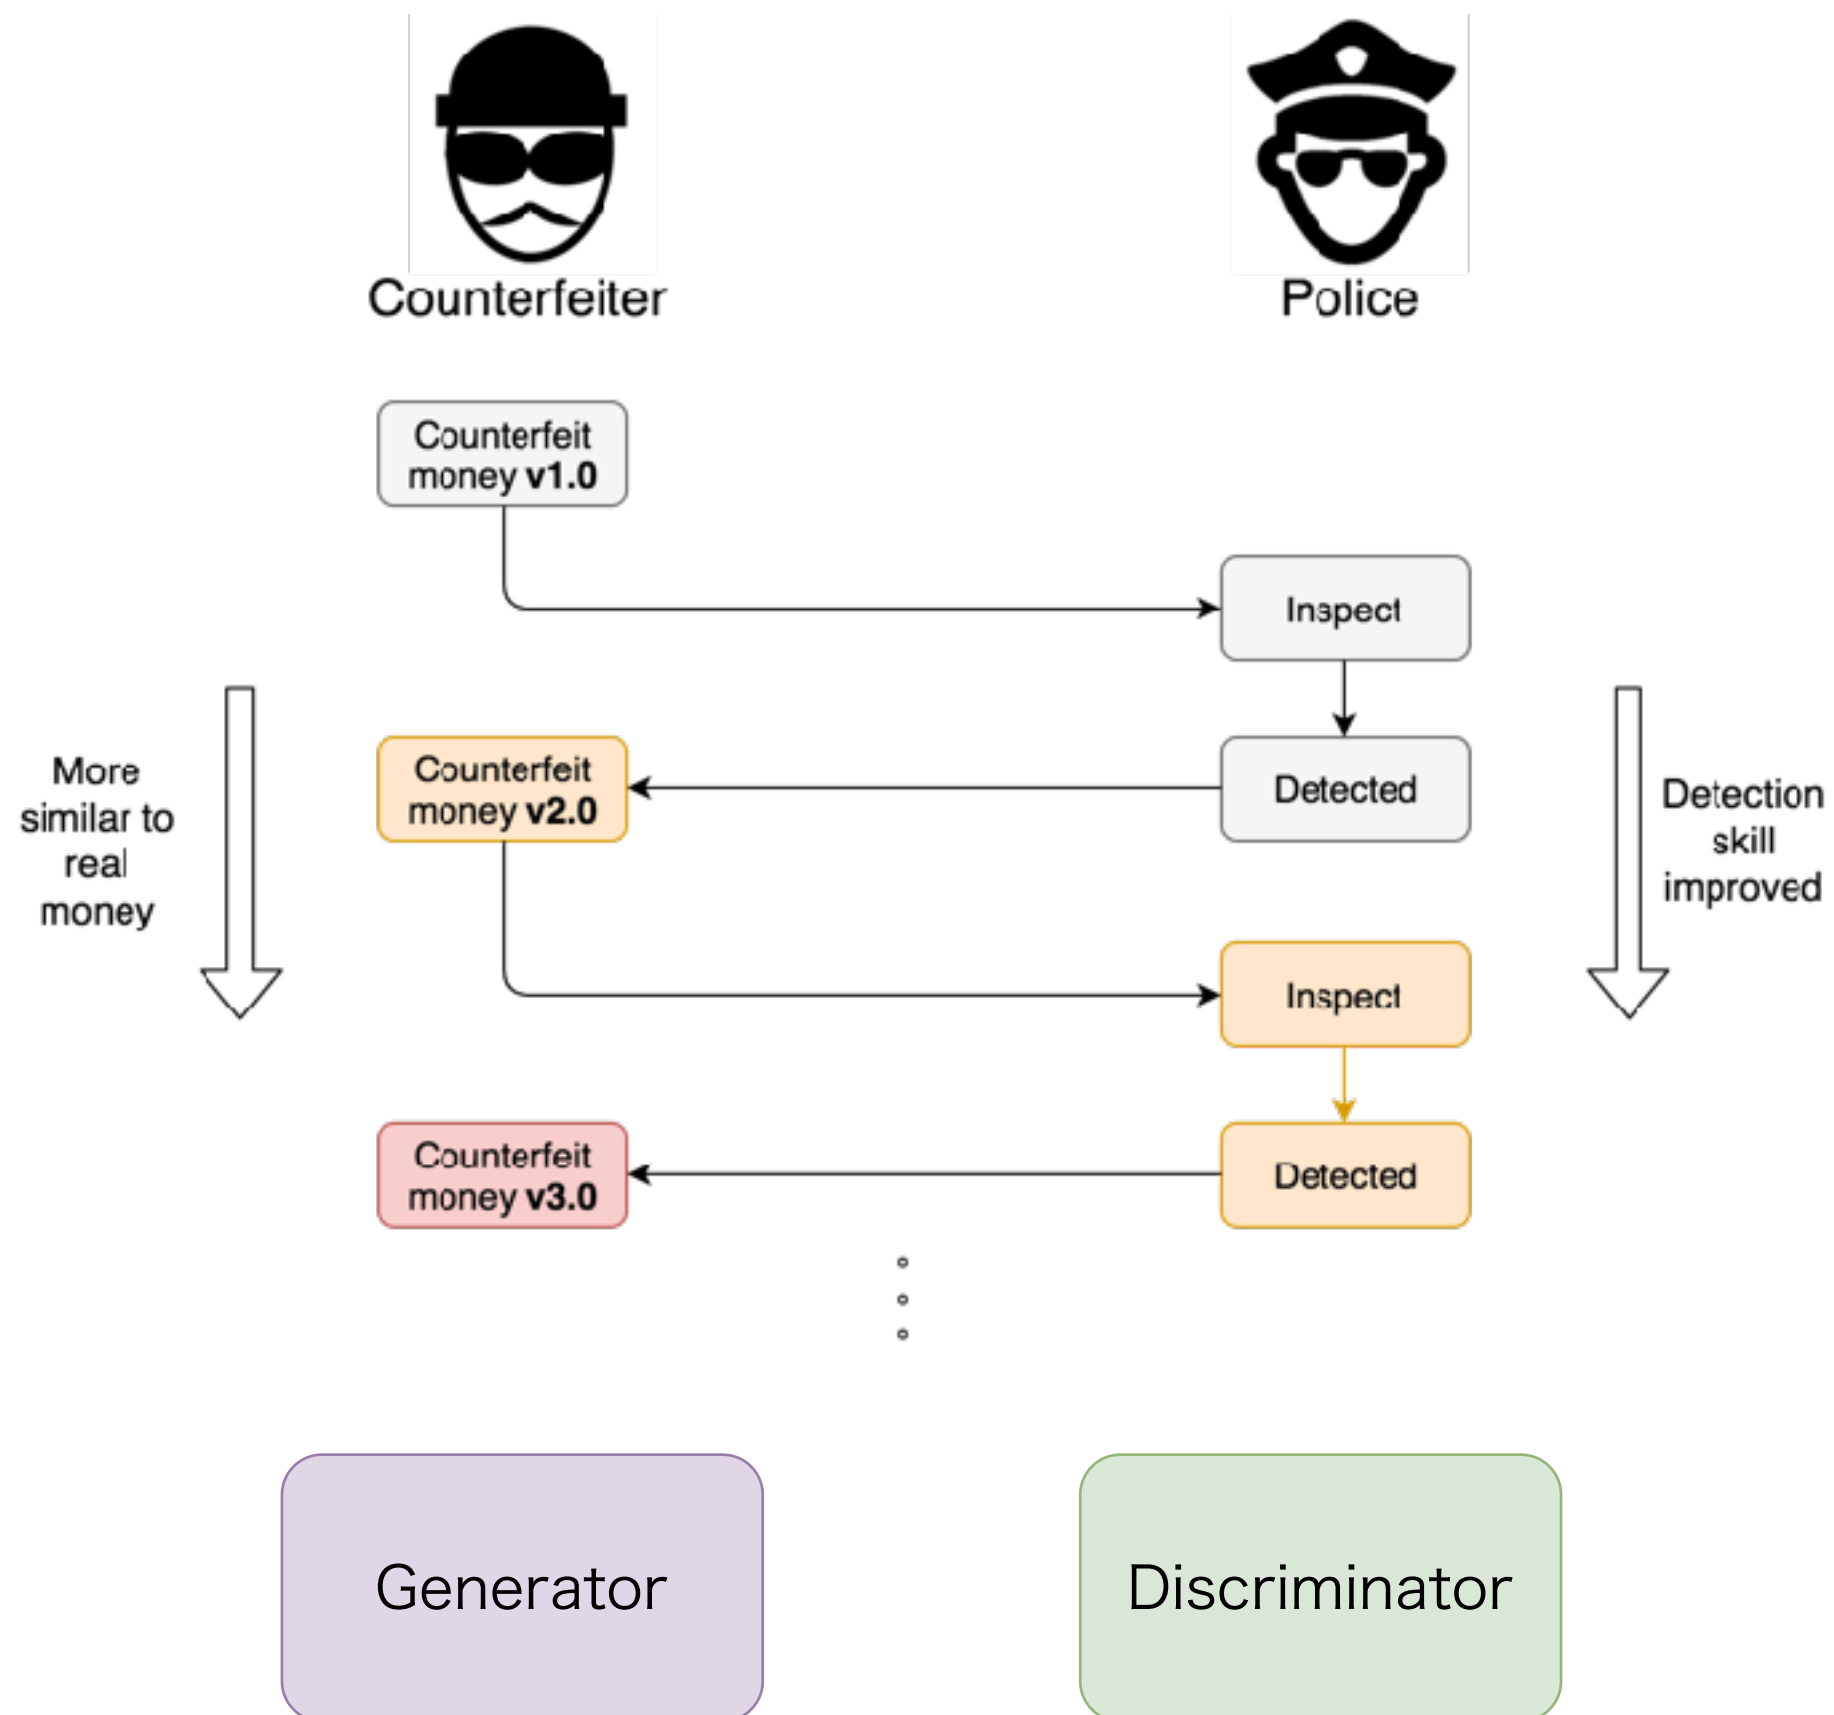


Model samples

(Goodfellow 2016)

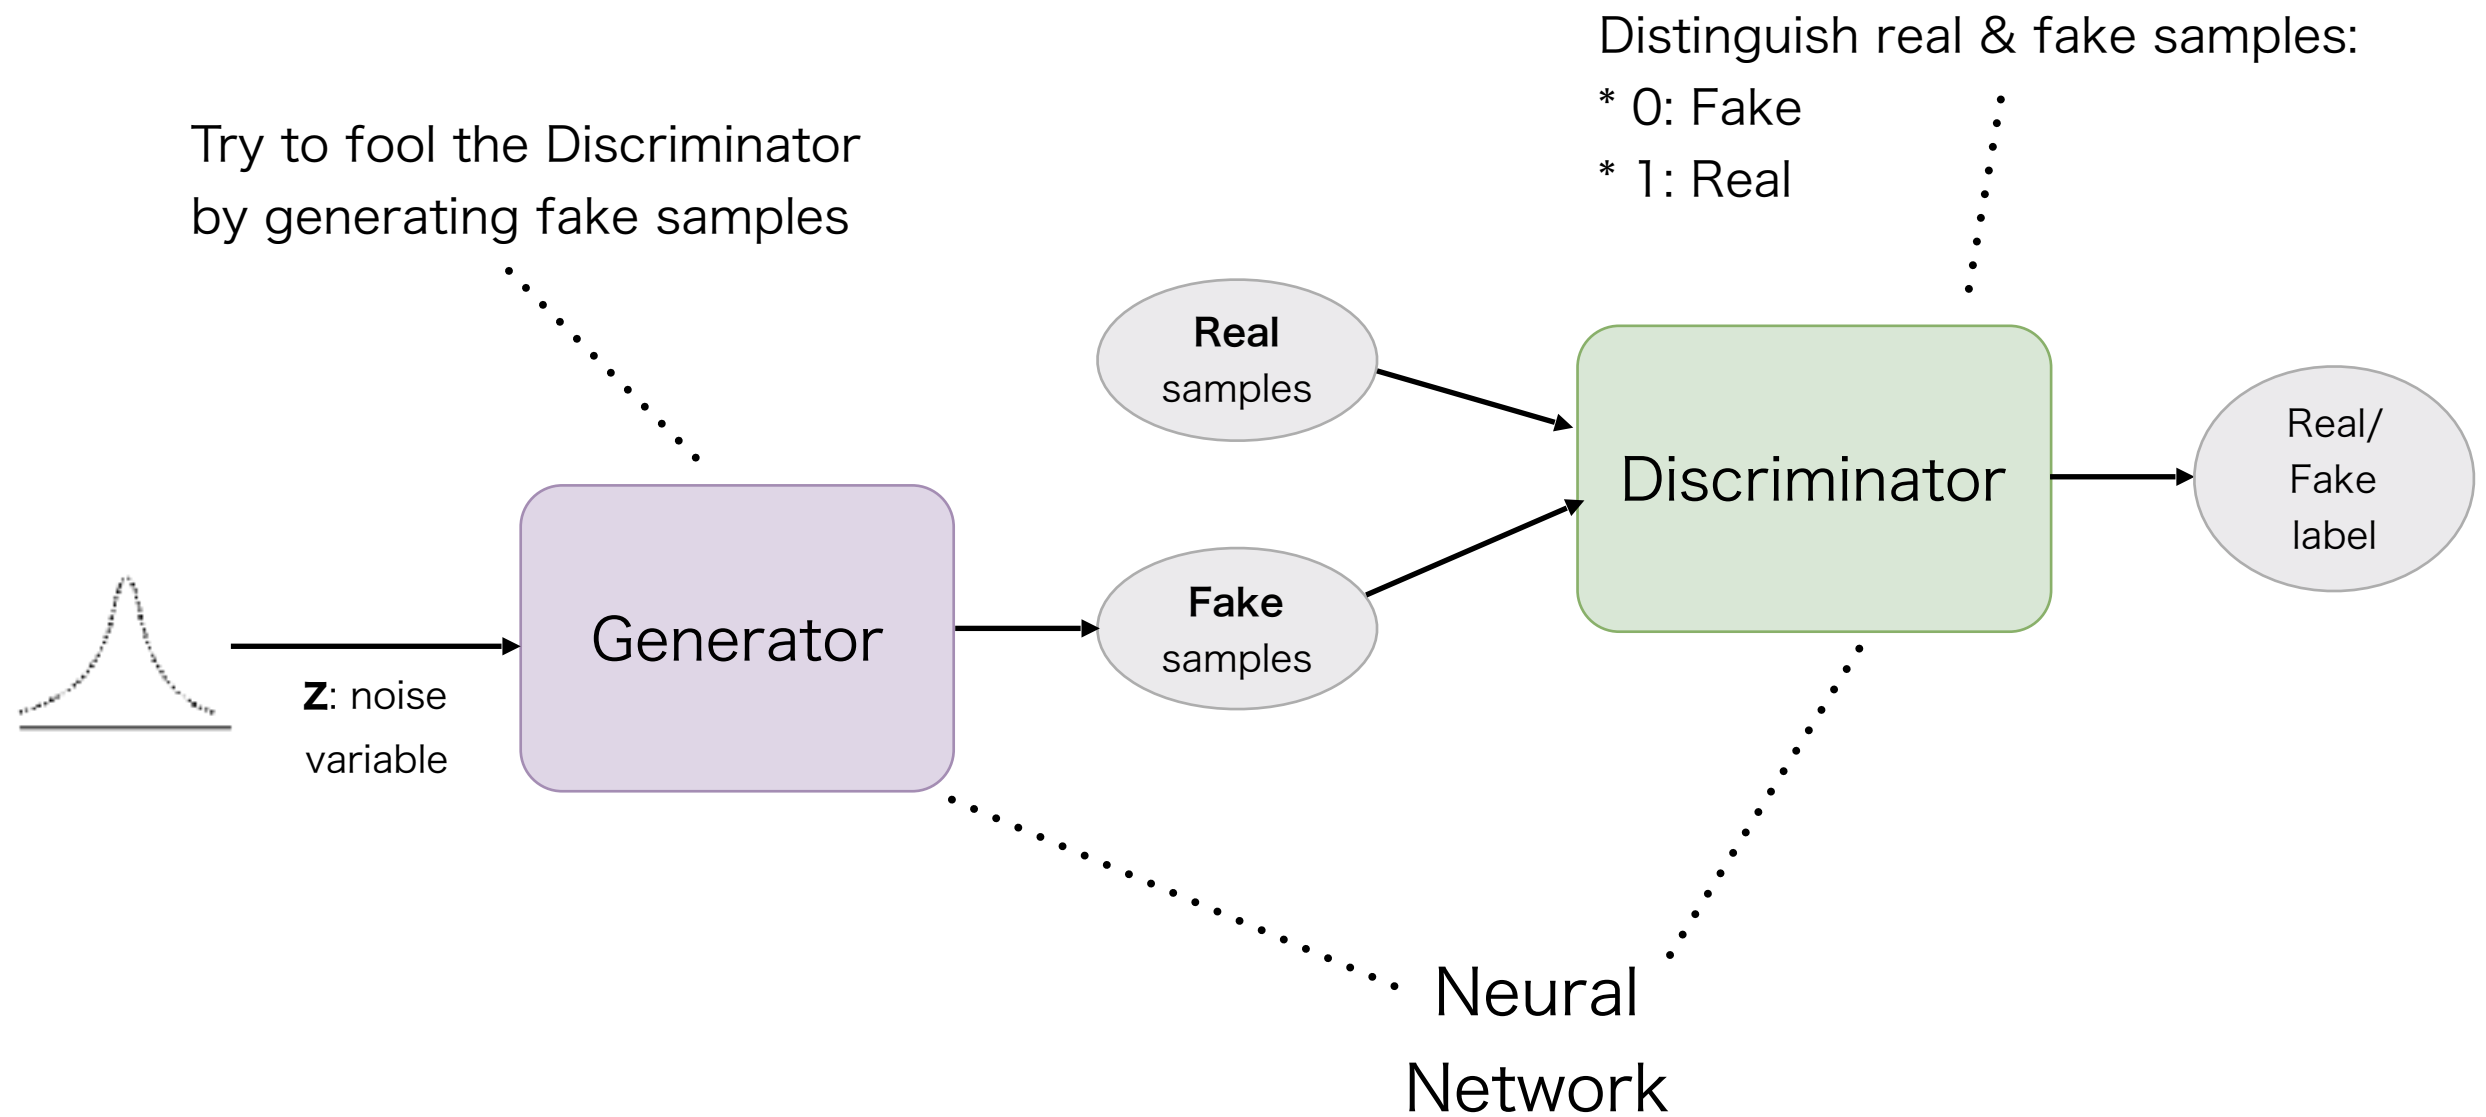
# The idea of GAN

# The idea of GAN



# Training GAN

# Framework



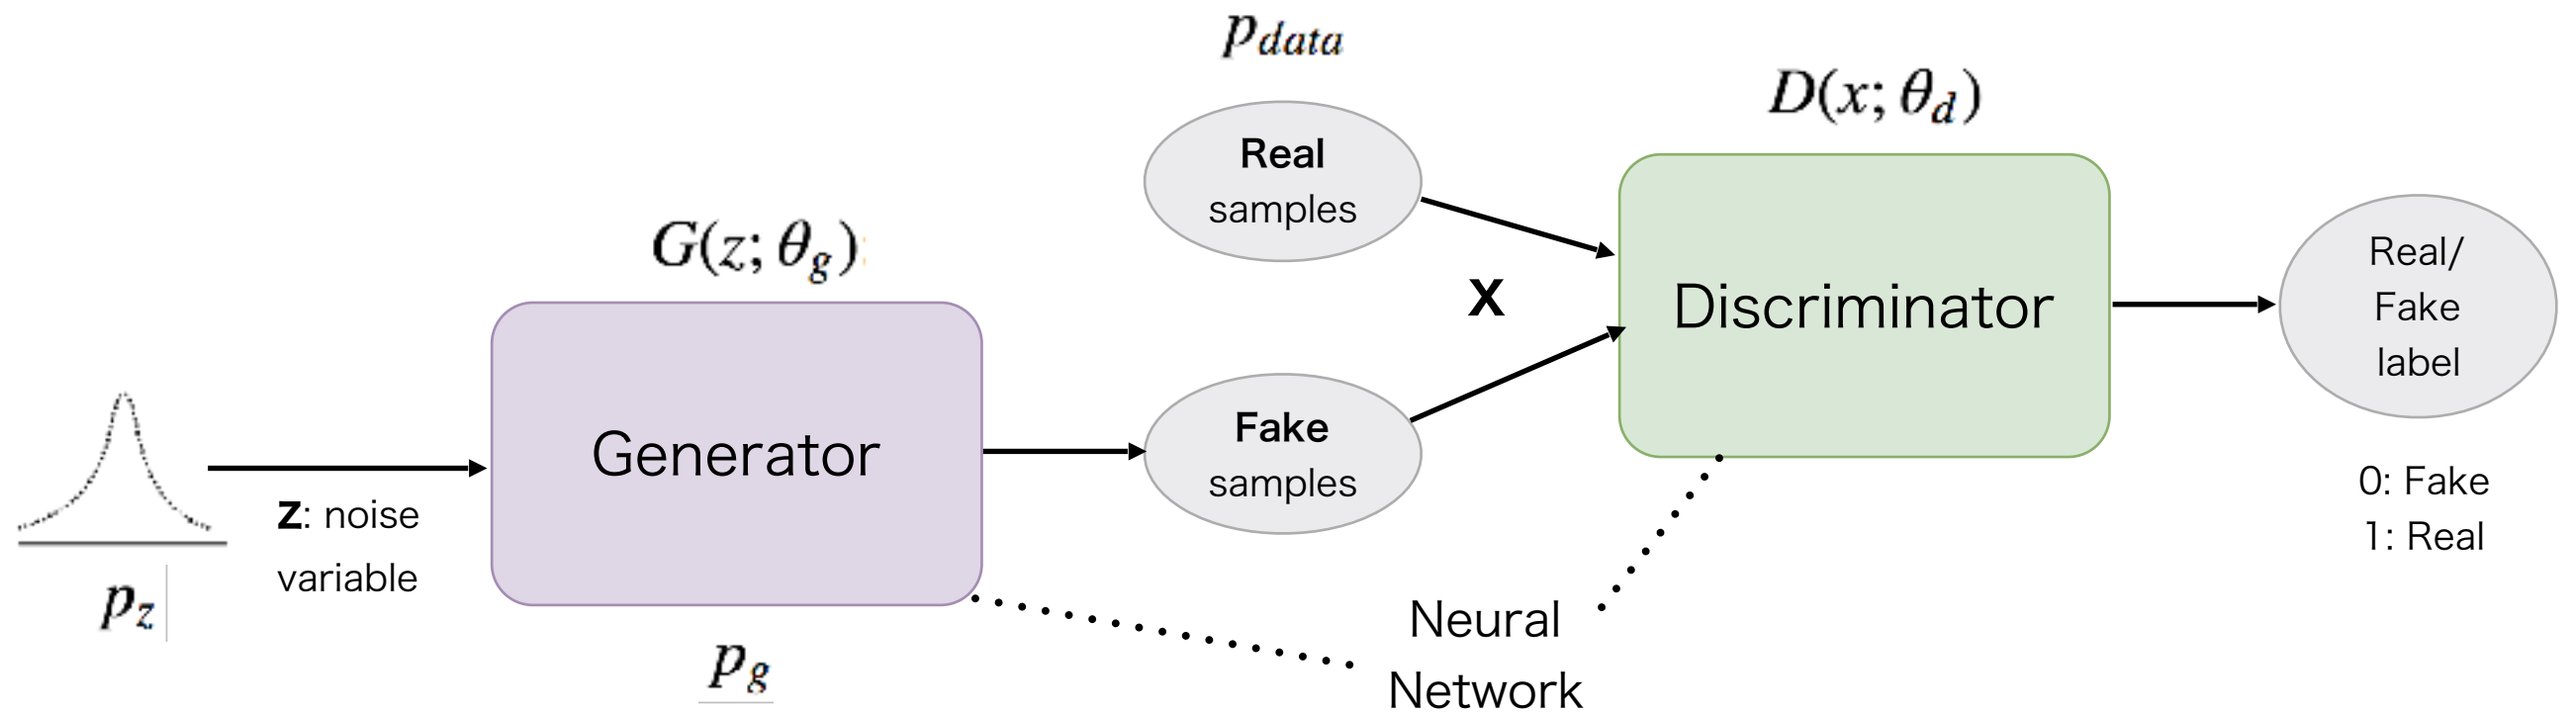
**Generative Adversarial Networks**



# Notation

$p_{data}$	Distribution of training data
$p_z$	Distribution of the noise variable $z$
$G(z; \theta_g)$	Differentiable function with parameters $\theta_g$ , which receives $z$ as input, and outputs a fake image.
$D(x; \theta_d)$	Function with parameters $\theta_d$ , which receives image $x$ as input, and outputs a probability that $x$ came from $p_{data}$
$p_g$	Generator's distribution
$m$	Number of training samples

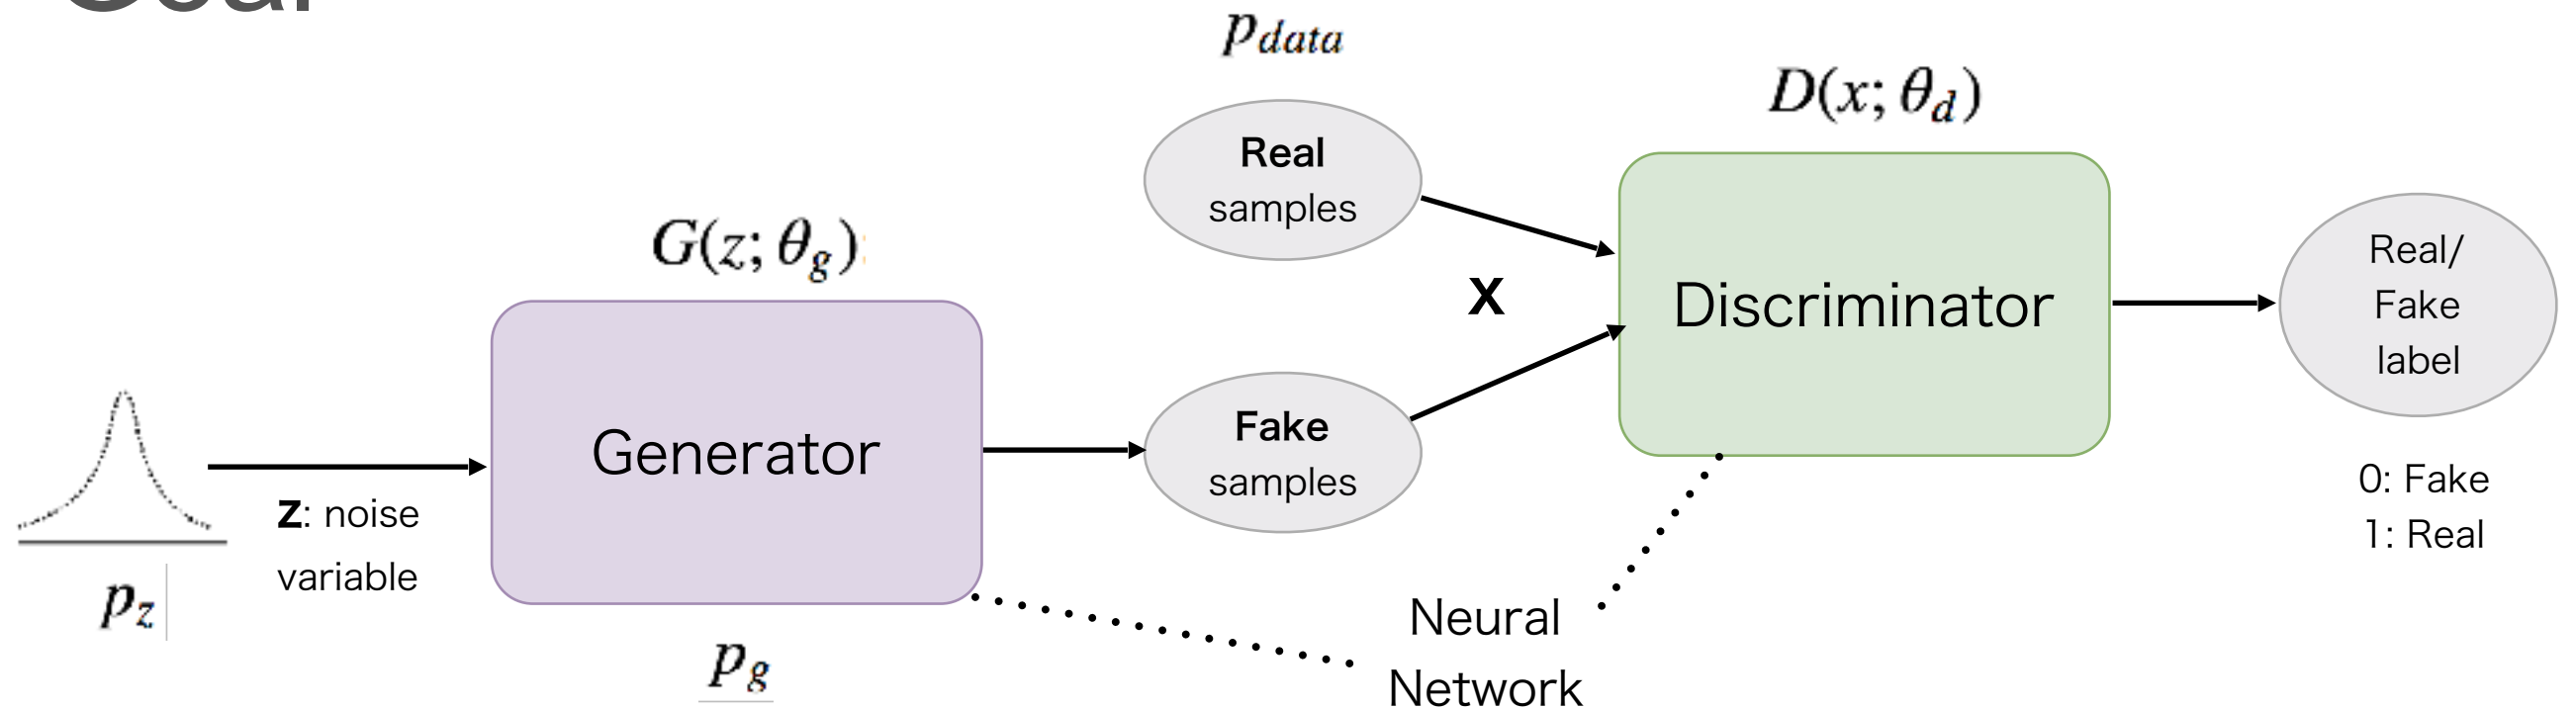
# Goal



## Ultimate goal:

- \* Train a good generator  $G$  which can generate an item  $x$  such that it seems to come from real samples.

# Goal



- Minimize the loss function of G

$$\min_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))$$

- Maximize the loss function of the D

$$\max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

# Minimax game

- The training process is playing the minimax game between G and D:

$$\min_{\theta_g} \max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

(Goodfellow 2014)

# Algorithm

---

**Algorithm 1** Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator,  $k$ , is a hyperparameter. We used  $k = 1$ , the least expensive option, in our experiments.

---

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
- Sample minibatch of  $m$  examples  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  from data generating distribution  $p_{\text{data}}(\mathbf{x})$ .
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(\mathbf{x}^{(i)}) + \log \left( 1 - D(G(\mathbf{z}^{(i)})) \right) \right].$$

**end for**

- Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log \left( 1 - D(G(\mathbf{z}^{(i)})) \right).$$


**end for**

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

---

# Gradient saturation in minimax game

- The inner loop (with k) will optimize D. We hope this step makes  $D(G(z))$  near to 0
- In minimax game, the loss function of G will be optimized with the target  $\log(1-D(G(z)))$ , which gives low gradient when  $D(G(z))$  near to 0



- One trick to solve the low gradient problem is to change the target when optimizing the loss function of G:

$$\min_{\theta_g} [-E_{z \sim p(z)} \log(D_{\theta_d}(G_{\theta_g}(z)))]$$

# Improved version

- Minimize the loss function of G

$$\min_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))$$

- Maximize the loss function of the D

$$\max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$



- Minimize the loss function of G

$$\min_{\theta_g} [-\mathbb{E}_{z \sim p(z)} \log(D_{\theta_d}(G_{\theta_g}(z)))]$$

- Maximize the loss function of the D

$$\max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

Improved version

(Goodfellow 2016)

# Distribution matching

(Jupyter notebook)




# Example



# Generating handwritten digits

# Example

- Generate hand-writing digit trained on MNIST dataset
- Train the models on MNIST dataset
  - Dataset of hand-written digit greyscale images
  - 60,000 training and 10,000 testing images of size 28x28



Real hand-writing digits  
(real images)



Generated hand-writing digits  
(fake images)

# Change the loss function optimization for practice

- Minimize the loss function of G

Improved version

$$\min_{\theta_g} [-E_{z \sim p(z)} \log(D_{\theta_d}(G_{\theta_g}(z)))]$$

- Maximize the loss function of the D

$$\max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$



- Minimize the loss function of G

In practice

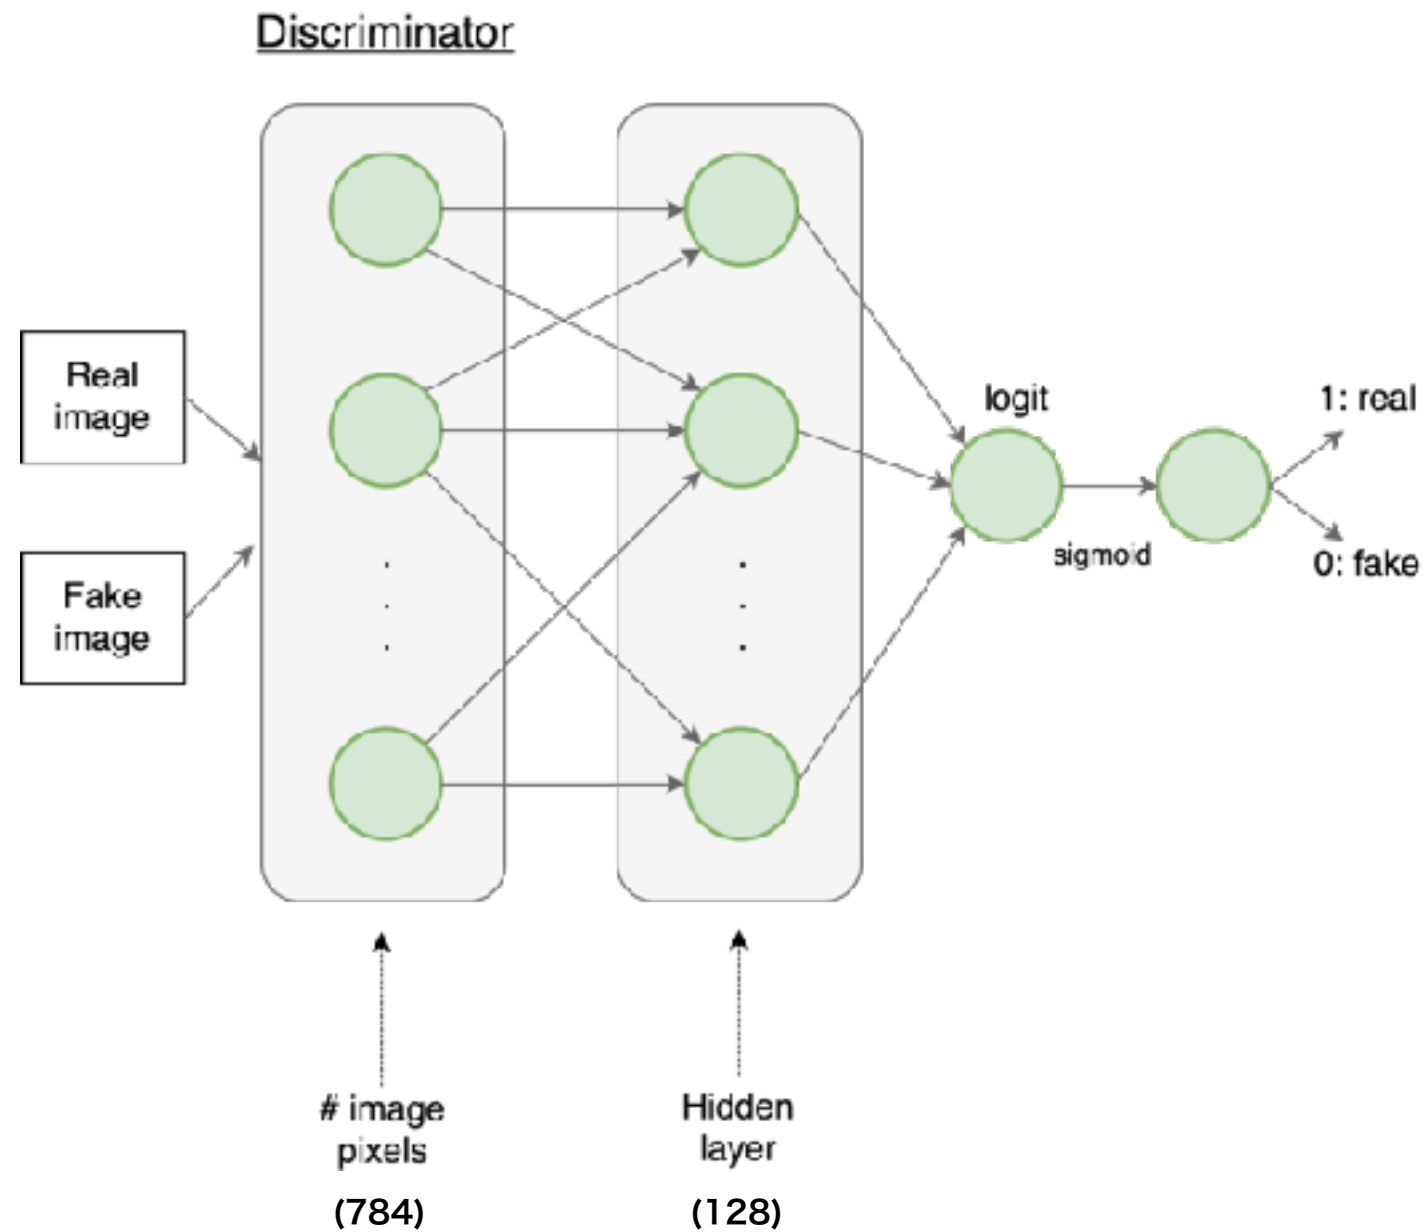
$$\min_{\theta_g} [-E_{z \sim p(z)} \log(D_{\theta_d}(G_{\theta_g}(z)))]$$

- Minimize the loss function of the D

$$\min_{\theta_d} [-E_{x \sim p_{data}} \log(D_{\theta_d}(x)) - E_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))]$$

(Goodfellow 2016)

# Network Architecture - Discriminator




# Network Architecture

$$G_{loss} = -\frac{1}{m} \sum_i (y_r \cdot \log \hat{y}_f + (1 - y_r) \cdot \log(1 - \hat{y}_f))$$


$$= -\frac{1}{m} \sum_i \log \hat{y}_f$$

Cross-Entropy loss  
btw predicted FAKE label  
& REAL label to fool D

Generator



Discriminator



$$D_{loss} = D_{loss\_real} + D_{loss\_fake}$$

Cross-Entropy loss

$$D_{loss\_real} = -\frac{1}{m} \sum_i (y_r \cdot \log \hat{y}_r + (1 - y_r) \cdot \log(1 - \hat{y}_r))$$

$$= -\frac{1}{m} \sum_i \log \hat{y}_r$$

$$D_{loss\_fake} = -\frac{1}{m} \sum_i (y_f \cdot \log \hat{y}_f + (1 - y_f) \cdot \log(1 - \hat{y}_f))$$

$$= -\frac{1}{m} \sum_i \log(1 - \hat{y}_f)$$


# Demo

(Jupyter notebook)

# Variants

# DCGAN

- DCGAN architecture for the generator used for Large-scale Scene Understanding (LSUN)
- From 100-dim random space to 64x64 images



(Radford et al, ICLR 2016)



# DCGAN

- Generated samples after training on LSUN dataset




(Radford et al, ICLR 2016)



# DCGAN

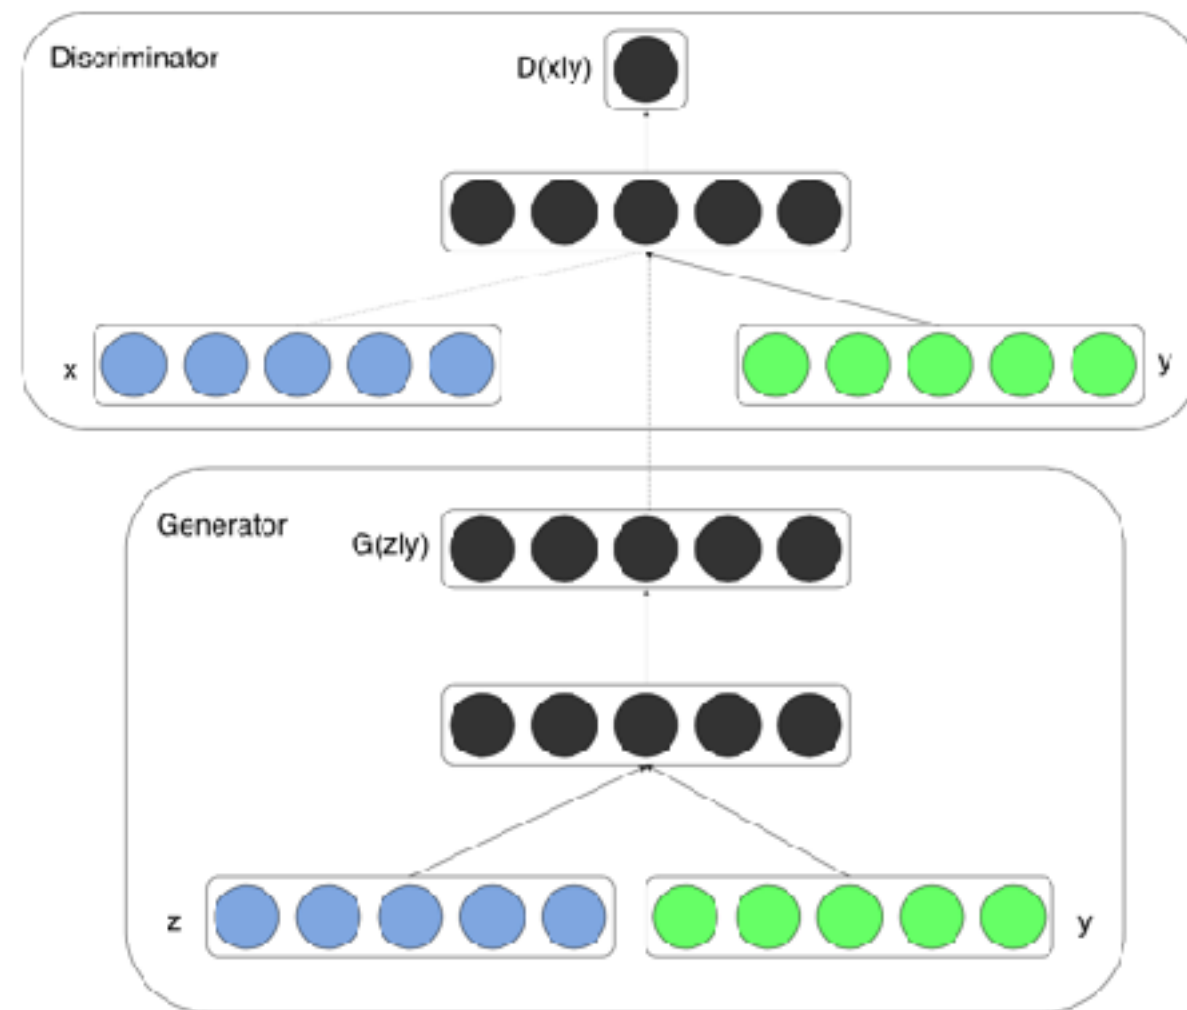
- Vector arithmetic for visual concepts



(Radford et al, ICLR 2016)

# Conditional GAN

- Class labels are encoded as one-hot vectors
- Control class labels of images by feeding one-hot vectors to both Discriminator & Generator



$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$




$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x}|\mathbf{y})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\log(1 - D(G(\mathbf{z}|\mathbf{y})))]$$

Conditions on  
the instances for both D & G



# GAN Zoo

- Near to 200 GAN variants since the vanilla GAN was proposed in 2014



- The GAN Zoo:
  - <https://deephunt.in/the-gan-zoo-79597dc8c347>

# Questions

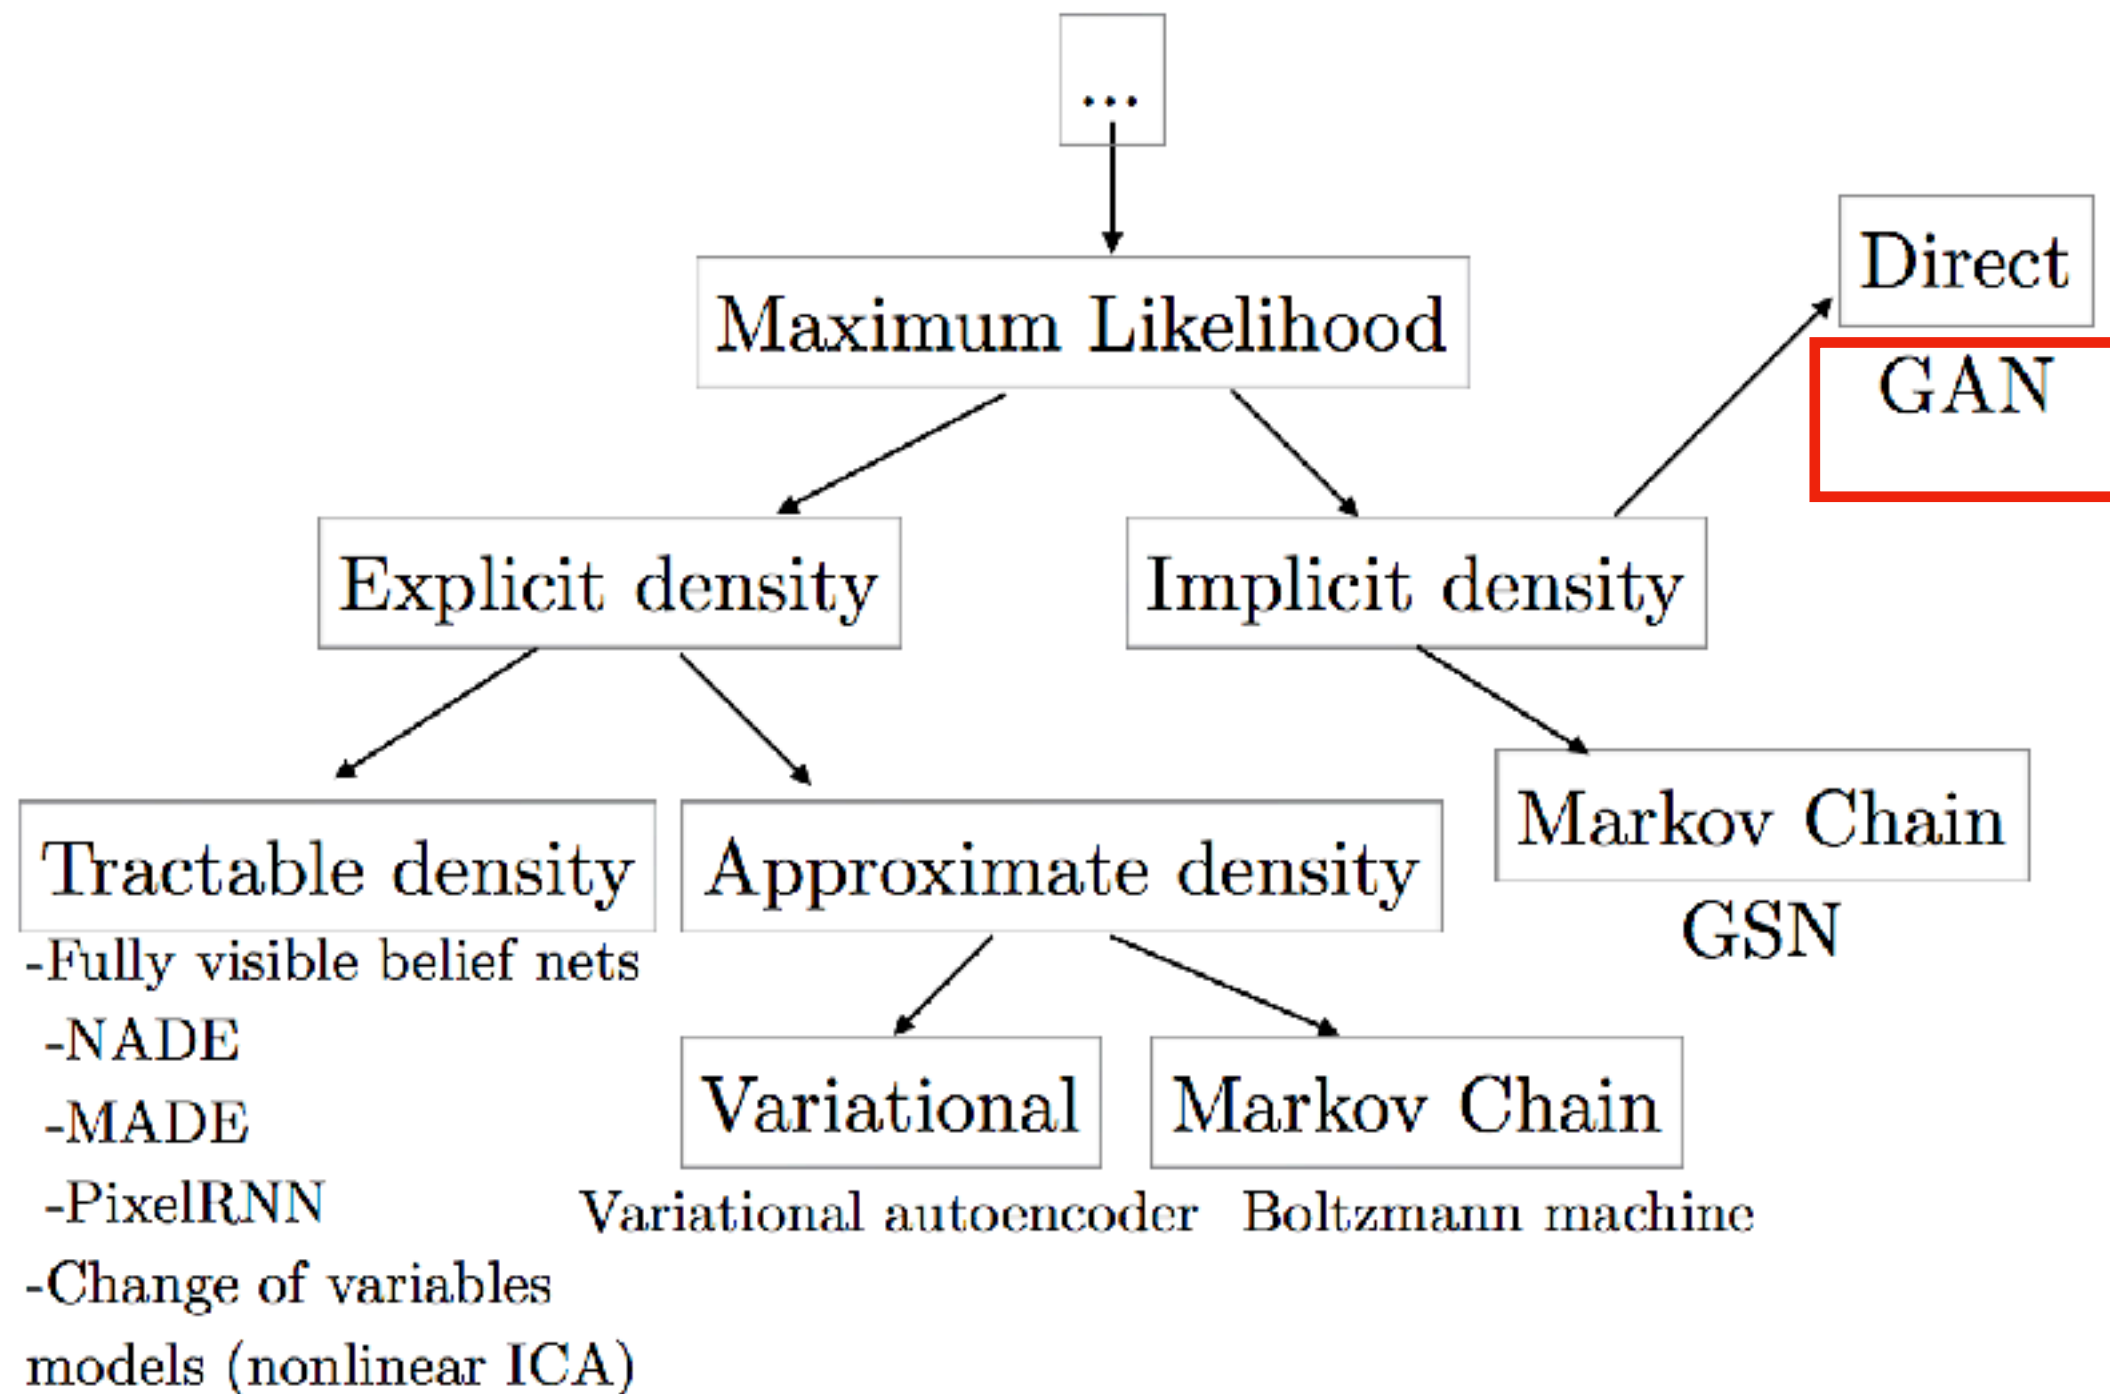
**Is there anyway to evaluate the  
generated images automatically?**  
(so hard with training  $\leftrightarrow$  looking)

# Applications of Generative Models

# Applications of Generative Models

- Create art works
- Build simulation environments
- Fix corrupted data...

# Generative Modeling methods



(Goodfellow 2016)



# Summary

- GAN trains a generator along with a discriminator by an adversarial process:
  - Discriminator aims to differentiate between real data and generated data
  - Generator tries to fool the discriminator by generating data
- Both of discriminator and generator are usually neural networks (although they can be other models)
- Proposed in 2014, GAN opened a new and hot research trend in Generative Modeling

# Thank you!

